

Lab3_js85773_asm3539

February 12, 2018

1 Problem 1

1.1 Part 1

We can consider $S = \text{span}\{v_1, v_2, v_3, v_4\}$. The explanations are printed out with the answers below

```
In [2]: import sympy
        from scipy import linalg
        import numpy as np
        import scipy

In [41]: # Setup Problem
        v1 = np.array([1, 2, 3, 4])
        v2 = np.array([0, 1, 0, 1])
        v3 = np.array([1, 4, 3, 6])
        v4 = np.array([2, 11, 6, 15])
        A = np.column_stack((v1, v2, v3, v4))

        # Create a vector inside S different from v1, v2, ...
        s = np.add(np.add(2 * v1, 2 * v2), np.add(2 * v3, 2 * v4))
        print("A vector in S that is not any of the given vectors is: {}".format(s))
        print("")

        # Create a vector not in S
        print("Any vector x with dim(x) > 4 is not in S")
        print("For example, x = [1, 1, 1, 1, 1]")
        print("")

        # Create a vector x perpendicular to S.
        # (any vector u s.t. v1 dot u, v2 dot u, ... etc are all equal to zero)
        A_transpose = sympy.Matrix(A.T)
        null_A = A_transpose.nullspace()
        null_A_ = np.array(null_A)
        #print(null_A_.T)
        print("Null(A^T) = {}".format(null_A_.T))
        # x = some vector orthogonal to S
        b = np.ones(2)
```

```

x = np.dot(null_A.T , b)
x = np.array(x, dtype="int")
print("A vector perpendicular to S is any vector in Null(A^T)")
print("So any linear combination, for instance: {}".format(x))

```

A vector in S that is not any of the given vectors is: [8 36 24 52]

Any vector x with $\dim(x) > 4$ is not in S
For example, $x = [1, 1, 1, 1, 1]$

```

Null(A^T) = [[-3 -2]
[0 -1]
[1 0]
[0 1]]
A vector perpendicular to S is any vector in Null(A^T)
So any linear combination, for instance: [-5 -1 1 1]

```

The answer to the question $u \in S$ for some new vector u is equivalent to asking if there is a solution to $Ax = u$. This can be solved using numpy or by hand with row reduction.

1.2 Part 2

Find the dimension of the subspace S.

```

In [35]: dim = np.linalg.matrix_rank(A)
print(dim)

```

2

1.3 Part 3

Construct an orthonormal basis. Notice only two vectors are required because the dimension of S is only 2.

```

In [36]: vectors = []
vectors.append(v1)
vectors.append(v3)
vectors.append(v2)
vectors.append(v4)
vectors = np.array(vectors)
#sympy.Matrix(vectors).rref()
vectors = vectors.T
orth = scipy.linalg.orth(vectors)
print(orth)
print(np.dot(orth[:,0], orth[:,1]))

```

```

[[-0.1098991  0.2669373 ]
 [-0.54396107 -0.53613402]
 [-0.32969731  0.80081189]
 [-0.76375927 -0.00225943]]
7.502679033599691e-17

```

1.4 Part 4

We are asked to solve

$$\min_{x \in S} \|\mathbf{x} - \mathbf{z}^*\|_2$$

$$\rightarrow \min_{x \in S} \sqrt{(\sum_{i=1}^n x_i - (1,0,0,0)^T)^2}$$

$$\rightarrow \min_{x \in S} (\sum_{i=1}^n x_i - (1,0,0,0)^T)^2$$

This is a convex function, so we set the derivative and set it equal to 0:

$$\sum_{i=1}^n (x_i - (1,0,0,0)^T) = 0$$

$$\rightarrow \sum_{i=1}^n x_i = (1,0,0,0)^T$$

Obviously we want $\mathbf{x} = \mathbf{z}^*$

Since $\mathbf{z}^* \notin S$, this becomes equivalent to finding the shortest distance to \mathbf{z}^*

We find this distance by projecting the vector \mathbf{z}^* onto S :

```

In [37]: min_dist = np.zeros(4)
         z_star = np.array([1, 0, 0, 0])
         col_1 = np.dot(orth[:,0], z_star)
         col_2 = np.dot(orth[:,1], z_star)
         min_dist = np.add(col_1, col_2)
         print("The minimum distance is: {}".format(min_dist))

```

The minimum distance is: 0.15703819557191717

```

In [38]: # Project u onto v
         #def proj(u, v):
         #    assert type(u) == np.ndarray
         #    assert type(v) == np.ndarray, "Invalid input"
         #    a = (np.dot(u, v))/(np.dot(u,u))
         #    w = a * v
         #    return w

         #z_star = sympy.Matrix(z_star)
         #space = sympy.Matrix(orth)
         #print(orth.shape[0])
         #print(z_star.T.shape)
         #space.project(z_star_)
         #min_dist = np.add(proj(z_star, orth[:,i]), min_dist)

```

1.5 Problem 2

1.5.1 Part 1 Most Common Words

```
In [118]: import requests as req
          from bs4 import BeautifulSoup
          from urllib.request import urlretrieve,urlopen
          import PyPDF2
          import io
          import pdfminer
          import os
          import numpy as np
          import nltk
          from math import log2
          from nltk.corpus import stopwords
          from nltk.tokenize import word_tokenize
          from io import StringIO
          from pdfminer.pdfinterp import PDFResourceManager, PDFPageInterpreter
          from pdfminer.converter import TextConverter
          from pdfminer.layout import LAParams
          from pdfminer.pdfpage import PDFPage
          import re
          from nltk.tokenize import RegexpTokenizer
          from collections import Counter

In [3]: #getting html from http://proceedings.mlr.press/v70/
        html_content = req.get("http://proceedings.mlr.press/v70/").content
        content = BeautifulSoup(html_content,'html.parser')

        #getting all the links
        links_pdf =[]

        #getting all the classes with id links
        links_paper = content.find_all(class_="links")

        #iterate over all the classes obtained and get link and keep only pdf links
        for link_paper in links_paper:
            links = link_paper.find_all('a')
            links_pdf.append(links[1].get('href'))

        i=0

        for link_pdf in links_pdf:
            content=urlopen(link_pdf).read()
            i = i+1
            filename = str(i)+".pdf"
            fout=open(filename, "wb")
            fout.write(content)
            fout.close()
```

```
print("downloaded {} files".format(i))
```

downloaded 434 files

```
In [15]: # function to convert text from pdf
def convert_pdf_to_txt(path):
    rsrcmgr = PDFResourceManager()
    retstr = io.StringIO()
    codec = 'utf-8'
    laparams = LAParams()
    device = TextConverter(rsrcmgr, retstr, codec=codec, laparams=laparams)
    fp = open(path, 'rb')
    interpreter = PDFPageInterpreter(rsrcmgr, device)
    password = ""
    maxpages = 0
    caching = True
    pagenos = set()

    for page in PDFPage.get_pages(fp, pagenos, maxpages=maxpages,
                                   password=password,
                                   caching=caching,
                                   check_extractable=True):
        interpreter.process_page(page)

    text = retstr.getvalue()

    fp.close()
    device.close()
    retstr.close()
    return text
```

```
In [73]: nltk.download('punkt')
         nltk.download('stopwords')
```

```
[nltk_data] Downloading package punkt to /home/humble-
[nltk_data]      fool/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /home/humble-
[nltk_data]      fool/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

```
Out[73]: True
```

```
In [75]: path = '/home/humble-fool/school_work/Spring_18/Data-Science-Lab-ee379k/lab3/src/'

         data_list = []
         for i in range(1,435):
```

```

# loads the pdf and sends it to the function ofr conversion
print(i)
t = convert_pdf_to_txt(path+str(i)+".pdf")
t_lower = t.lower()
## tokenizes and removes punctuation
tokenizer = RegexpTokenizer(r'\w+')
tokens = tokenizer.tokenize(t_lower)

#another filter to remove stop words
filtered_tokens = tokens[:]
for token in tokens:
    if token in stopwords.words('english'):
        filtered_tokens.remove(token)
#another filter to remove unnecessary element of length lesser than three

count_filter = []
for token in filtered_tokens:
    if (len(token) >=3):
        count_filter.append(token)
data_list.append(count_filter)

```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26

27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74

75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122

123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170

171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201

KeyboardInterrupt

Traceback (most recent call last)

```
<ipython-input-75-394b78917016> in <module>()
      5     # loads the pdf and sends it to the function ofr conversion
      6     print(i)
----> 7     t = convert_pdf_to_txt(path+str(i)+".pdf")
      8     t_lower = t.lower()
      9     ## tokenizes and removes punctuation
```

```
<ipython-input-15-e97870ead6f1> in convert_pdf_to_txt(path)
     16                                     caching=caching,
```

```

17                                     check_extractable=True):
---> 18         interpreter.process_page(page)
19
20         text = retstr.getvalue()

~/school_work/Spring_18/Data-Science-Lab-ee379k/lib/python3.5/site-packages/pdfminer/pdf
851         self.device.begin_page(page, ctm)
852         self.render_contents(page.resources, page.contents, ctm=ctm)
--> 853         self.device.end_page(page)
854         return
855

~/school_work/Spring_18/Data-Science-Lab-ee379k/lib/python3.5/site-packages/pdfminer/con
52         assert isinstance(self.cur_item, LTPage), str(type(self.cur_item))
53         if self.laparams is not None:
---> 54             self.cur_item.analyze(self.laparams)
55             self.pageno += 1
56             self.receive_layout(self.cur_item)

~/school_work/Spring_18/Data-Science-Lab-ee379k/lib/python3.5/site-packages/pdfminer/lay
684         textboxes = list(self.group_textlines(laparams, textlines))
685         if -1 <= laparams.bboxes_flow and laparams.bboxes_flow <= +1 and textboxes:
--> 686             self.groups = self.group_textboxes(laparams, textboxes)
687             assigner = IndexAssigner()
688             for group in self.groups:

~/school_work/Spring_18/Data-Science-Lab-ee379k/lib/python3.5/site-packages/pdfminer/lay
661             plane.remove(obj1)
662             plane.remove(obj2)
--> 663             dists = [ (c,d,obj1,obj2) for (c,d,obj1,obj2) in dists
664                         if (obj1 in plane and obj2 in plane) ]
665             for other in plane:

~/school_work/Spring_18/Data-Science-Lab-ee379k/lib/python3.5/site-packages/pdfminer/lay
662             plane.remove(obj2)
663             dists = [ (c,d,obj1,obj2) for (c,d,obj1,obj2) in dists
--> 664                         if (obj1 in plane and obj2 in plane) ]
665             for other in plane:
666                 dists.append((0, dist(group, other), group, other))

```

KeyboardInterrupt:

```
In [80]: print(len(data_list))
        print(len(data_list[50]))
```

```
200
2888
```

```
In [82]: ## pdf number 201 didn't load for some reason
```

```
path = '/home/humble-fool/school_work/Spring_18/Data-Science-Lab-ee379k/lab3/src/'

for i in range(202,435):
    # loads the pdf and sends it to the function ofr conversion
    print(i)
    t = convert_pdf_to_txt(path+str(i)+".pdf")
    t_lower = t.lower()
    ## tokenizes and removes punctuation
    tokenizer = RegexpTokenizer(r'\w+')
    tokens = tokenizer.tokenize(t_lower)

    #another filter to remove stop words
    filtered_tokens = tokens[:]
    for token in tokens:
        if token in stopwords.words('english'):
            filtered_tokens.remove(token)
    #another filter to remove unnecessary element of length lesser than three

    count_filter = []
    for token in filtered_tokens:
        if (len(token) >=3):
            count_filter.append(token)
    data_list.append(count_filter)
```

```
202
203
204
205
206
207
208
209
210
211
212
213
214
215
```

216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263

264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311

312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359

360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407

408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434

```
In [84]: print(len(data_list))
```

433

```
In [87]: one_list = []  
         for data in data_list:  
             one_list = one_list+data  
         print(len(one_list))
```

1648762

```
In [91]: most_common_words= [word for word, word_count in Counter(one_list).most_common(10)]  
         print (most_common_words)
```

```
['cid', 'learning', 'model', 'data', 'algorithm', '2016', '107', 'set', 'function', '2015']
```

1.5.2 Part 2 Estimate the Entropy

```
In [125]: ## randomly choosing a list corresponding to an ICML paper
          rand_list = np.random.choice(data_list)
          rand_word = np.random.choice(rand_list)
          freq = 0
          for word in rand_list:
              if word == rand_word:
                  freq += 1

          prob = freq/len(rand_list)
          entropy = -(prob*log2(prob)+(1-prob)*log2(1-prob))

          print(entropy)

          ## need to figure out the entropy formula they want

0.0012351778656126482
0.013714018232625748
```

1.5.3 Part 3 Random Paragraph

```
In [105]: unique_words = set(one_list)
          corresponding_prob = []

          for word in unique_words:
              #calculate number of occurrences
              num_occurrences = one_list.count(word)
              #calc prob
              probability = num_occurrences/len(one_list)
              corresponding_prob.append(probability)

In [104]: print(len(corresponding_prob))

131

In [113]: a = np.random.choice(one_list, 500, corresponding_prob)
          para = ""
          for word in a:
              para = para + word
          para += " "

In [114]: print(para)
```

lan sition ieee increase riemannian iot note nfsic abilistic matrix tion getting ference values

1.6 Problem 3


1.6.1 Part 2 Following the tutorial

```
In [52]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib
import os
from sklearn.linear_model import Ridge
from sklearn.model_selection import cross_val_score
import matplotlib.pyplot as plt
from scipy.stats import skew
from scipy.stats.stats import pearsonr
from sklearn.metrics import mean_squared_error
from math import sqrt

%config InlineBackend.figure_format = 'retina' #set 'png' here when working on notebook
%matplotlib inline

In [38]: train = pd.read_csv("train.csv")
test = pd.read_csv("test.csv")
all_data = pd.concat((train.loc[:, 'MSSubClass': 'SaleCondition'], test.loc[:, 'MSSubClass': 'SaleCondition']))

In [39]: matplotlib.rcParams['figure.figsize'] = (15.0, 6.0)
prices = pd.DataFrame({'price': train["SalePrice"], "log(price+1)": np.log1p(train["SalePrice"])})
_ = prices.hist()
```



```
In [40]: train['SalePrice'] = np.log1p(train["SalePrice"])
numeric_feats = all_data.dtypes[(all_data.dtypes != 'object')].index
skewed_feats = train[numeric_feats].apply(lambda x: skew(x.dropna()))
skewed_feats = skewed_feats[skewed_feats>0.75]
```

```

skewed_feats = skewed_feats.index
all_data[skewed_feats] = np.log1p(all_data[skewed_feats])
all_data = pd.get_dummies(all_data)
all_data = all_data.fillna(all_data.mean())
X_train = all_data[:train.shape[0]]
X_test = all_data[train.shape[0]:]
y_train = train.SalePrice

```

```

In [41]: model = Ridge(alpha = 0.1)
         model.fit(X_train,y_train)

```

```

Out[41]: Ridge(alpha=0.1, copy_X=True, fit_intercept=True, max_iter=None,
              normalize=False, random_state=None, solver='auto', tol=0.001)

```

```

In [42]: arr = np.expm1(model.predict(X_test))
         pd.DataFrame(arr).to_csv("kaggle_results.csv")

```

```

In [43]: rmse_train = sqrt(mean_squared_error(model.predict(X_train), y_train))
         print("The training RMSE is: {}".format(rmse_train))

```

The training RMSE is: 0.09211955585640522

1.6.2 Part 3 Trying a different model

```

In [72]: # Naive model- find best training error

```


```

X_train = np.array(X_train)
y_train = np.array(y_train)
alphas = np.arange(0.1, 1, 0.025)
rmse_train = []

for alpha in alphas:
    model = Ridge(alpha)
    model.fit(X_train, y_train)
    rmse_train = sqrt(mean_squared_error(model.predict(X_train), y_train))
    rmse_train.append(rmse_train)

_ = plt.scatter(alphas, rmse_train)
_ = plt.title("Training Error for Varying Alphas")
_ = plt.ylabel("RMSE")
_ = plt.xlabel("Alpha")

```



The alpha we chose originally has the lowest training error. Without cross-validation, this is the best ridge regression we can choose.